CS 6476 Project 4

Mengying Lin
mlin365@gatech.edu
mlin365
904016037

Part 1: SimpleNet

[Insert loss plot for SimpleNet here]

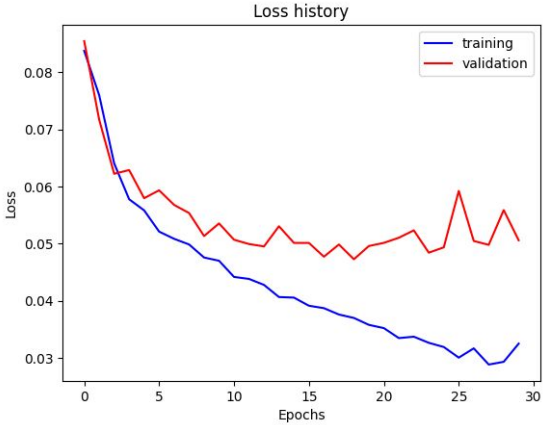

[Insert accuracy plot for SimpleNet here]

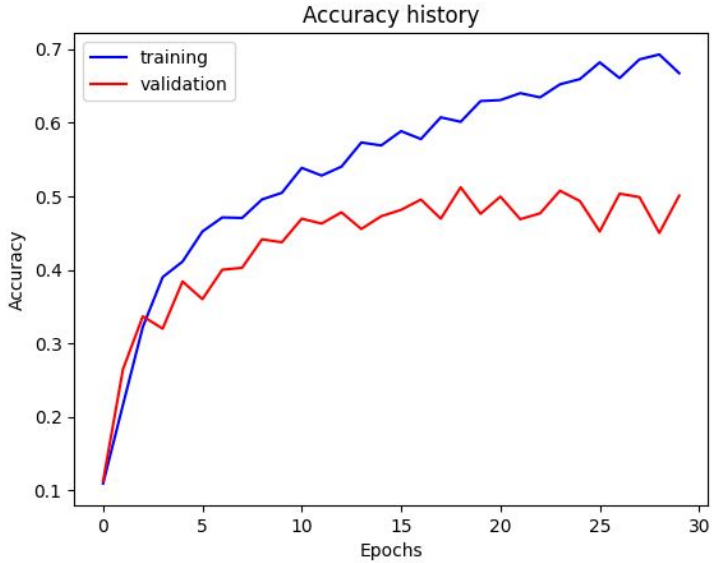

Final training accuracy:

Final validation accuracy:

`Train Accuracy = 0.6670016750618442; Validation Accuracy = 0.5006666668256123`

Part 2: SimpleNetFinal

Add each of the following (keeping the changes as you move to the next row):

	Training accuracy	Validation accuracy
SimpleNet	0.67	0.50
+ Jittering	0.65	0.55
+ Zero-centering & variance-normalization	0.72	0.65
+ Dropout regularization	0.70	0.72
+ Making network "deep"	0.83	0.81
+ Batch normalization	0.89	0.90

Part 2: SimpleNetFinal

[Insert loss plot for SimpleNetFinal here]

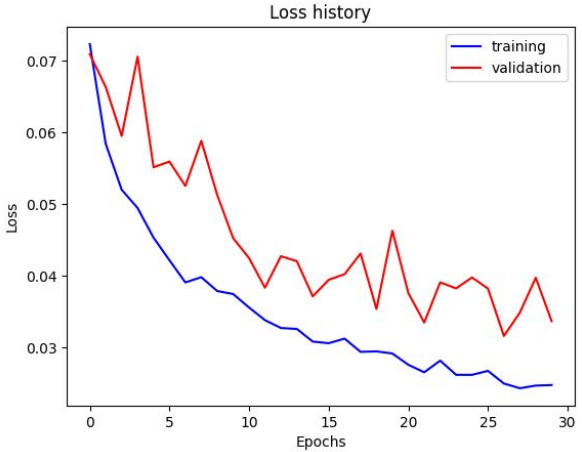

[Insert accuracy plot for SimpleNetFinal here]

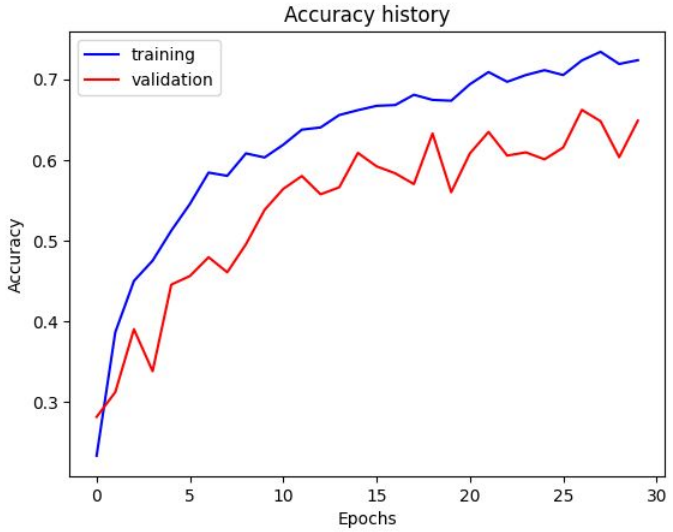

Final training accuracy:

Final validation accuracy:

Train Accuracy = 0.7236180904921974;

Validation Accuracy = 0.648666666507721

Part 2: SimpleNetFinal

[Name 10 different possible transformations for data augmentation.]
Horizontal Flip, Vertical Flip, Rotation, Scaling, Translation, Random Crop, Color Jitter, Gaussian Noise, Brightness Adjustment, Shear Transformation

[What is the desired variance after each layer? Why would that be helpful?]
The desired variance after each layer in a neural network is 1. This is helpful because it ensures that the inputs to each layer have a consistent scale, which prevents exploding or vanishing gradients and stabilizes learning across layers.

Part 2: SimpleNetFinal

[What distribution is dropout usually sampled from?]

Bernoulli distribution

[How many parameters does your base SimpleNet model have? How many parameters does your SimpleNetFinal model have?]

Simplenet: 56,895
SimpleNetFinal: 126,383

[What is the effect of batch norm after a conv layer with a bias compared to having no batch norm?]

With Batch Norm: The effect of the bias in the convolutional layer is negated because Batch Norm subtracts the mean and normalizes the activations, making the bias unnecessary.
Without Batch Norm: The bias is used to shift the activations before they are passed to the next layer.

Part 3: ResNet

[Insert loss plot here]

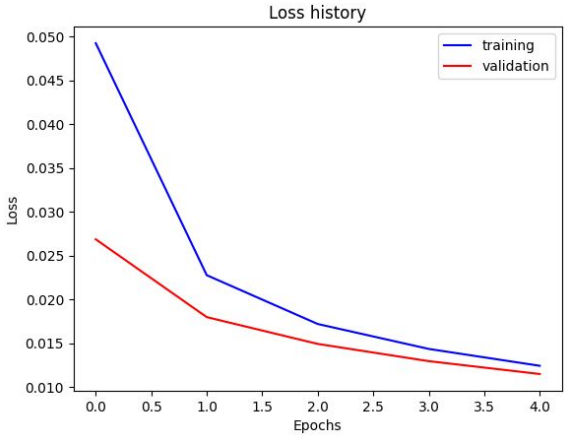

[Insert accuracy plot here]

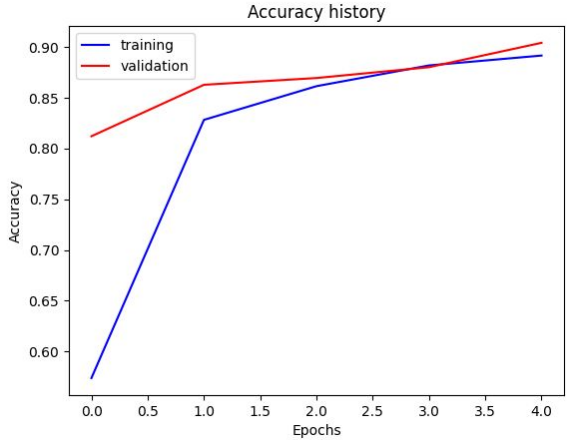

Final training accuracy:

Final validation accuracy:

Train Accuracy = 0.8914572864321608; Validation Accuracy = 0.9040000003178914

Part 3: ResNet

[Insert visualization of confusion matrix obtained from your final ResNet model.]

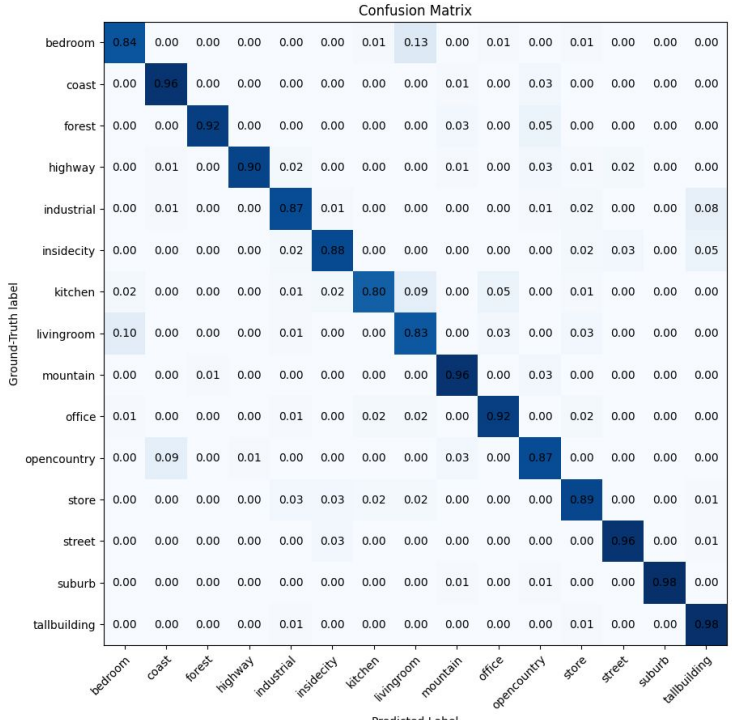

Part 3: ResNet

[Insert visualizations of 3 misclassified images from the most misclassified class according to your confusion matrix. Explain why this may have occurred.]

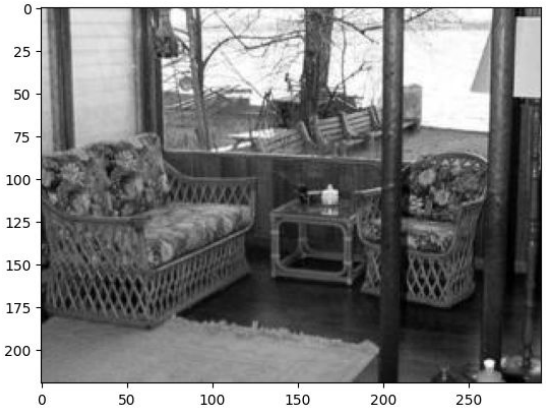

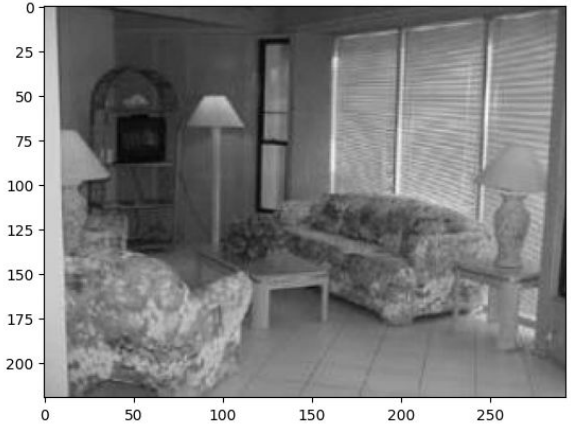

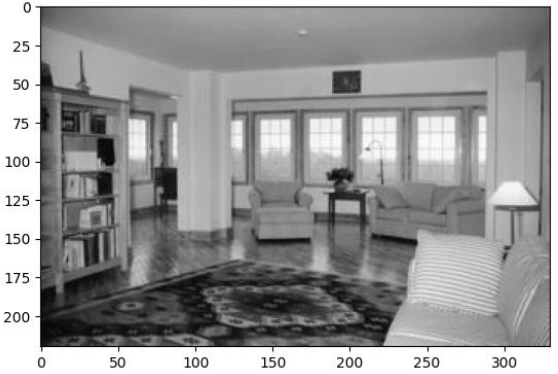

All of the above are labeled as ‘bedroom,’ while their ground truth (GT) labels are ‘living room.’ This might be because they all have a sofa in the scene, which looks similar to a bed, leading to confusion.

Part 3: ResNet

[What does fine-tuning a network mean?]

It means adjust the model architecture (e.g. adding more layers, adding normalizations, adding dropout, adjusting the number of neurons within layer) to make the performance better.

[Why do we want to "freeze" the conv layers and some of the linear layers from a pre-trained ResNet? Why can we do this?]

The pretrained model serves as a feature extractor, and we are typically unwilling to let the initially “bad” gradients to affect the well-established feature extractor. We can do this because the gradients won’t flow back to those layers if their parameters are detached from the computational graph.

Part 4: Logging to Weights and Biases

[Add the link to the weights and Biases project below. Make sure your project is **public.**]

View run **resnet** at: https://wandb.ai/cassielin0910/cv-assn4/runs/ltqvob9k

View project at: https://wandb.ai/cassielin0910/cv-assn4

Extra credit

[Add the link to the Weights & Biases run for the model used in the “huggingface.py” script below]

View run of Mobilevit-small (worse than ResNet) at https://wandb.ai/cassielin0910/cv-assn4/runs/8q3tuasx
View run of ConvNext (better than ResNet) at https://wandb.ai/cassielin0910/cv-assn4/runs/66jkfvbi